

Complex Wishart sampling

Ray Hinton

Expectations

In the following two examples, the sample mean and true mean are compared when two different methods are used to generate complex Wishart matrices.

First, the method using a product of matrices is shown. The true mean and sample mean are close in Frobenius distance, and their entries appear similar.

```
# Frobenius distance
norm(df*Sigma - avgW, "F")
```

```
[1] 0.9593366
```

```
# true mean
round(df*Sigma, 3)
```

| | [,1] | [,2] | [,3] | [,4] |
|------|-----------------|-----------------|-----------------|-----------------|
| [1,] | 51.918+ 0.000i | -11.404-16.387i | -32.573-37.767i | 16.797+ 7.337i |
| [2,] | -11.404+16.387i | 14.281+ 0.000i | 24.559+10.457i | -2.703+ 1.385i |
| [3,] | -32.573+37.767i | 24.559-10.457i | 173.866+ 0.000i | -19.671-30.567i |
| [4,] | 16.797- 7.337i | -2.703- 1.385i | -19.671+30.567i | 19.827+ 0.000i |

```
# sample mean
round(avgW, 3)
```

| | [,1] | [,2] | [,3] | [,4] |
|------|-----------------|-----------------|-----------------|-----------------|
| [1,] | 51.849+ 0.000i | -11.389-16.387i | -32.461-37.573i | 16.800+ 7.319i |
| [2,] | -11.389+16.387i | 14.280+ 0.000i | 24.482+10.410i | -2.710+ 1.399i |
| [3,] | -32.461+37.573i | 24.482-10.410i | 173.001+ 0.000i | -19.628-30.421i |
| [4,] | 16.800- 7.319i | -2.710- 1.399i | -19.628+30.421i | 19.778+ 0.000i |

Next, the method based on the Cholesky decomposition is shown. Again, the true means and sample means are similar.

```
# Frobenius distance
norm(df*Sigma - avgW, "F")
```

```
[1] 0.5050177
```

```
# true mean
round(df*Sigma, 3)
```

```
      [,1]      [,2]      [,3]      [,4]
[1,] 51.918+ 0.000i -11.404-16.387i -32.573-37.767i 16.797+ 7.337i
[2,] -11.404+16.387i 14.281+ 0.000i 24.559+10.457i -2.703+ 1.385i
[3,] -32.573+37.767i 24.559-10.457i 173.866+ 0.000i -19.671-30.567i
[4,] 16.797- 7.337i -2.703- 1.385i -19.671+30.567i 19.827+ 0.000i
```

```
# sample mean
round(avgW, 3)
```

```
      [,1]      [,2]      [,3]      [,4]
[1,] 52.055+ 0.000i -11.452-16.411i -32.754-37.697i 16.931+ 7.375i
[2,] -11.452+16.411i 14.307+ 0.000i 24.564+10.455i -2.722+ 1.407i
[3,] -32.754+37.697i 24.564-10.455i 174.058+ 0.000i -19.824-30.655i
[4,] 16.931- 7.375i -2.722- 1.407i -19.824+30.655i 19.950+ 0.000i
```

Speed

There are two parts of the algorithm that can be slow. First, the matrix square root of the Σ parameter is needed. This can be computed just using the eigendecomposition via `eigen`. Alternatively, `EigenR::Eigen_sqrt` is faster in some cases. Second, sampling the Cholesky decomposition of the complex Wishart matrix is generally faster as the degrees of freedom increases.

For small matrices and low degrees of freedom, finding $\Sigma^{1/2}$ with `EigenR::Eigen_sqrt` greatly speeds up the sampler. Using the Cholesky decomposition sampling technique does not make as much of a difference.

The following times are in microseconds for $p = 8$.

```

df <- 100
p <- 8

Sigma <- rcomplex_wishart(p + 1, diag(p))

mb_res <- microbenchmark(
  rcomplex_wishart(df, Sigma, useEigenR = FALSE, byCholesky = FALSE),
  rcomplex_wishart(df, Sigma, useEigenR = FALSE, byCholesky = TRUE),
  rcomplex_wishart(df, Sigma, useEigenR = TRUE, byCholesky = FALSE),
  rcomplex_wishart(df, Sigma, useEigenR = TRUE, byCholesky = TRUE),
  times = 1000
) |> summary()

mb_res$expr <- c("FALSE, FALSE", "FALSE, TRUE", "TRUE, FALSE", "TRUE, TRUE")
names(mb_res)[1] <- c("EigenR, byChol")

cat(paste0("df = ", df, ", p = ", p))

```

df = 100, p = 8

```
knitr::kable(mb_res, digits = 3)
```

| EigenR, byChol | min | lq | mean | median | uq | max | neval |
|----------------|---------|---------|---------|---------|---------|----------|-------|
| FALSE, FALSE | 338.928 | 416.659 | 648.890 | 511.708 | 774.944 | 4642.590 | 1000 |
| FALSE, TRUE | 306.442 | 387.097 | 611.252 | 481.089 | 740.568 | 5735.023 | 1000 |
| TRUE, FALSE | 80.678 | 90.098 | 136.938 | 101.998 | 168.229 | 685.329 | 1000 |
| TRUE, TRUE | 42.682 | 56.654 | 101.319 | 71.406 | 131.447 | 2360.553 | 1000 |

With a larger degrees of freedom, $df = 1000$, using the Cholesky decomposition also results in a speed up, regardless of how the matrix square root is found.

df = 1000, p = 8

| EigenR, byChol | min | lq | mean | median | uq | max | neval |
|----------------|---------|---------|----------|---------|----------|----------|-------|
| FALSE, FALSE | 783.915 | 870.781 | 1175.097 | 953.114 | 1303.448 | 4581.723 | 1000 |
| FALSE, TRUE | 309.284 | 397.941 | 545.834 | 455.643 | 629.077 | 1809.689 | 1000 |
| TRUE, FALSE | 517.809 | 538.770 | 741.213 | 558.704 | 798.628 | 6263.800 | 1000 |

| EigenR, byChol | min | lq | mean | median | uq | max | neval |
|----------------|--------|--------|--------|--------|---------|----------|-------|
| TRUE, TRUE | 42.937 | 58.036 | 90.606 | 71.321 | 100.697 | 1654.859 | 1000 |

With larger matrices, $p = 64$ below, the `EigenR::Eigen_sqrt` function is slower than using `eigen` to compute the matrix square root. Using the Cholesky decomposition is faster in either case.

The following are in milliseconds.

$df = 100, p = 64$

| EigenR, byChol | min | lq | mean | median | uq | max | neval |
|----------------|-------|-------|-------|--------|-------|--------|-------|
| FALSE, FALSE | 2.611 | 2.738 | 3.503 | 2.982 | 3.994 | 8.456 | 100 |
| FALSE, TRUE | 2.443 | 2.531 | 3.115 | 2.645 | 3.359 | 7.280 | 100 |
| TRUE, FALSE | 3.530 | 3.585 | 4.139 | 3.652 | 3.908 | 10.599 | 100 |
| TRUE, TRUE | 3.319 | 3.429 | 3.837 | 3.466 | 3.674 | 7.572 | 100 |

$df = 1000, p = 64$

| EigenR, byChol | min | lq | mean | median | uq | max | neval |
|----------------|-------|-------|--------|--------|--------|--------|-------|
| FALSE, FALSE | 6.425 | 6.653 | 9.528 | 7.870 | 11.292 | 53.597 | 100 |
| FALSE, TRUE | 2.420 | 2.614 | 3.162 | 2.809 | 3.554 | 5.796 | 100 |
| TRUE, FALSE | 7.310 | 7.497 | 10.056 | 8.892 | 11.550 | 22.147 | 100 |
| TRUE, TRUE | 3.444 | 3.542 | 4.135 | 3.627 | 4.423 | 7.657 | 100 |

With even higher degrees of freedom, the Cholesky decomposition is clearly faster.

$df = 10000, p = 64$

| EigenR, byChol | min | lq | mean | median | uq | max | neval |
|----------------|--------|--------|--------|--------|--------|--------|-------|
| FALSE, FALSE | 47.935 | 52.195 | 59.754 | 56.873 | 63.654 | 95.190 | 100 |
| FALSE, TRUE | 2.493 | 3.004 | 3.921 | 3.407 | 4.330 | 12.022 | 100 |
| TRUE, FALSE | 48.263 | 51.104 | 59.340 | 56.608 | 63.326 | 97.877 | 100 |
| TRUE, TRUE | 3.398 | 3.750 | 5.009 | 3.964 | 4.307 | 40.335 | 100 |